

MINI PROJECT

1.Problem Statement:Which model is suitable best for Insurance Dataset

```
In [2]: 1 import pandas as pd
        2 import numpy as np
        3 from sklearn import preprocessing,svm
        4 import matplotlib.pyplot as plt
        5 import seaborn as sns
        6 from sklearn.model_selection import train_test_split
        7 from sklearn.linear_model import LinearRegression
        8 from sklearn.linear_model import Ridge
        9 from sklearn.linear_model import RidgeCV
       10 from sklearn.linear_model import Lasso
       11 from sklearn.linear_model import LassoCV
       12 from sklearn.linear_model import ElasticNet
       13 from sklearn import metrics
```

Data Collection

Read the Data

```
In [3]: 1 df=pd.read_csv(r"C:\Users\DELL E5490\Downloads\insurance.csv")
        2 df
```

```
Out[3]:
```

| | age | sex | bmi | children | smoker | region | charges |
|------|-----|--------|--------|----------|--------|-----------|-------------|
| 0 | 19 | female | 27.900 | 0 | yes | southwest | 16884.92400 |
| 1 | 18 | male | 33.770 | 1 | no | southeast | 1725.55230 |
| 2 | 28 | male | 33.000 | 3 | no | southeast | 4449.46200 |
| 3 | 33 | male | 22.705 | 0 | no | northwest | 21984.47061 |
| 4 | 32 | male | 28.880 | 0 | no | northwest | 3866.85520 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 1333 | 50 | male | 30.970 | 3 | no | northwest | 10600.54830 |
| 1334 | 18 | female | 31.920 | 0 | no | northeast | 2205.98080 |
| 1335 | 18 | female | 36.850 | 0 | no | southeast | 1629.83350 |
| 1336 | 21 | female | 25.800 | 0 | no | southwest | 2007.94500 |
| 1337 | 61 | female | 29.070 | 0 | yes | northwest | 29141.36030 |

1338 rows × 7 columns

2.Data Cleaning and Preprocessing

In [4]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

In [5]:

```
1 df.columns
```

Out[5]: Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'], dtype='object')

In [6]:

```
1 df.head()
```

Out[6]:

| | age | sex | bmi | children | smoker | region | charges |
|---|-----|--------|--------|----------|--------|-----------|-------------|
| 0 | 19 | female | 27.900 | 0 | yes | southwest | 16884.92400 |
| 1 | 18 | male | 33.770 | 1 | no | southeast | 1725.55230 |
| 2 | 28 | male | 33.000 | 3 | no | southeast | 4449.46200 |
| 3 | 33 | male | 22.705 | 0 | no | northwest | 21984.47061 |
| 4 | 32 | male | 28.880 | 0 | no | northwest | 3866.85520 |

In [7]:

```
1 df.tail()
```

Out[7]:

| | age | sex | bmi | children | smoker | region | charges |
|------|-----|--------|-------|----------|--------|-----------|------------|
| 1333 | 50 | male | 30.97 | 3 | no | northwest | 10600.5483 |
| 1334 | 18 | female | 31.92 | 0 | no | northeast | 2205.9808 |
| 1335 | 18 | female | 36.85 | 0 | no | southeast | 1629.8335 |
| 1336 | 21 | female | 25.80 | 0 | no | southwest | 2007.9450 |
| 1337 | 61 | female | 29.07 | 0 | yes | northwest | 29141.3603 |

In [8]:

```
1 df.shape
```

Out[8]: (1338, 7)

In [9]:

```
1 df.describe()
```

Out[9]:

| | age | bmi | children | charges |
|-------|-------------|-------------|-------------|--------------|
| count | 1338.000000 | 1338.000000 | 1338.000000 | 1338.000000 |
| mean | 39.207025 | 30.663397 | 1.094918 | 13270.422265 |
| std | 14.049960 | 6.098187 | 1.205493 | 12110.011237 |
| min | 18.000000 | 15.960000 | 0.000000 | 1121.873900 |
| 25% | 27.000000 | 26.296250 | 0.000000 | 4740.287150 |
| 50% | 39.000000 | 30.400000 | 1.000000 | 9382.033000 |
| 75% | 51.000000 | 34.693750 | 2.000000 | 16639.912515 |
| max | 64.000000 | 53.130000 | 5.000000 | 63770.428010 |

To find Duplicate Values

```
In [10]: 1 df.duplicated().sum()
```

```
Out[10]: 1
```

To find Unique Values

```
In [11]: 1 df['age'].unique()  
2 df['children'].unique()  
3 df['bmi'].unique()  
4
```

```

Out[11]: array([27.9 , 33.77 , 33.   , 22.705, 28.88 , 25.74 , 33.44 , 27.74 ,
29.83 , 25.84 , 26.22 , 26.29 , 34.4  , 39.82 , 42.13 , 24.6  ,
30.78 , 23.845, 40.3  , 35.3  , 36.005, 32.4  , 34.1  , 31.92 ,
28.025, 27.72 , 23.085, 32.775, 17.385, 36.3  , 35.6  , 26.315,
28.6  , 28.31 , 36.4  , 20.425, 32.965, 20.8  , 36.67 , 39.9  ,
26.6  , 36.63 , 21.78 , 30.8  , 37.05 , 37.3  , 38.665, 34.77 ,
24.53 , 35.2  , 35.625, 33.63 , 28.   , 34.43 , 28.69 , 36.955,
31.825, 31.68 , 22.88 , 37.335, 27.36 , 33.66 , 24.7  , 25.935,
22.42 , 28.9  , 39.1  , 36.19 , 23.98 , 24.75 , 28.5  , 28.1  ,
32.01 , 27.4  , 34.01 , 29.59 , 35.53 , 39.805, 26.885, 38.285,
37.62 , 41.23 , 34.8  , 22.895, 31.16 , 27.2  , 26.98 , 39.49 ,
24.795, 31.3  , 38.28 , 19.95 , 19.3  , 31.6  , 25.46 , 30.115,
29.92 , 27.5  , 28.4  , 30.875, 27.94 , 35.09 , 29.7  , 35.72 ,
32.205, 28.595, 49.06 , 27.17 , 23.37 , 37.1  , 23.75 , 28.975,
31.35 , 33.915, 28.785, 28.3  , 37.4  , 17.765, 34.7  , 26.505,
22.04 , 35.9  , 25.555, 28.05 , 25.175, 31.9  , 36.   , 32.49 ,
25.3  , 29.735, 38.83 , 30.495, 37.73 , 37.43 , 24.13 , 37.145,
39.52 , 24.42 , 27.83 , 36.85 , 39.6  , 29.8  , 29.64 , 28.215,
37.   , 33.155, 18.905, 41.47 , 30.3  , 15.96 , 33.345, 37.7  ,
27.835, 29.2  , 26.41 , 30.69 , 41.895, 30.9  , 32.2  , 32.11 ,
31.57 , 26.2  , 30.59 , 32.8  , 18.05 , 39.33 , 32.23 , 24.035,
36.08 , 22.3  , 26.4  , 31.8  , 26.73 , 23.1  , 23.21 , 33.7  ,
33.25 , 24.64 , 33.88 , 38.06 , 41.91 , 31.635, 36.195, 17.8  ,
24.51 , 22.22 , 38.39 , 29.07 , 22.135, 26.8  , 30.02 , 35.86 ,
20.9  , 17.29 , 34.21 , 25.365, 40.15 , 24.415, 25.2  , 26.84 ,
24.32 , 42.35 , 19.8  , 32.395, 30.2  , 29.37 , 34.2  , 27.455,
27.55 , 20.615, 24.3  , 31.79 , 21.56 , 28.12 , 40.565, 27.645,
31.2  , 26.62 , 48.07 , 36.765, 33.4  , 45.54 , 28.82 , 22.99 ,
27.7  , 25.41 , 34.39 , 22.61 , 37.51 , 38.   , 33.33 , 34.865,
33.06 , 35.97 , 31.4  , 25.27 , 40.945, 34.105, 36.48 , 33.8  ,
36.7  , 36.385, 34.5  , 32.3  , 27.6  , 29.26 , 35.75 , 23.18 ,
25.6  , 35.245, 43.89 , 20.79 , 30.5  , 21.7  , 21.89 , 24.985,
32.015, 30.4  , 21.09 , 22.23 , 32.9  , 24.89 , 31.46 , 17.955,
30.685, 43.34 , 39.05 , 30.21 , 31.445, 19.855, 31.02 , 38.17 ,
20.6  , 47.52 , 20.4  , 38.38 , 24.31 , 23.6  , 21.12 , 30.03 ,
17.48 , 20.235, 17.195, 23.9  , 35.15 , 35.64 , 22.6  , 39.16 ,
27.265, 29.165, 16.815, 33.1  , 26.9  , 33.11 , 31.73 , 46.75 ,
29.45 , 32.68 , 33.5  , 43.01 , 36.52 , 26.695, 25.65 , 29.6  ,
38.6  , 23.4  , 46.53 , 30.14 , 30.   , 38.095, 28.38 , 28.7  ,
33.82 , 24.09 , 32.67 , 25.1  , 32.56 , 41.325, 39.5  , 34.3  ,
31.065, 21.47 , 25.08 , 43.4  , 25.7  , 27.93 , 39.2  , 26.03 ,

```


```

30.25 , 28.93 , 35.7 , 35.31 , 31. , 44.22 , 26.07 , 25.8 ,
39.425, 40.48 , 38.9 , 47.41 , 35.435, 46.7 , 46.2 , 21.4 ,
23.8 , 44.77 , 32.12 , 29.1 , 37.29 , 43.12 , 36.86 , 34.295,
23.465, 45.43 , 23.65 , 20.7 , 28.27 , 35.91 , 29. , 19.57 ,
31.13 , 21.85 , 40.26 , 33.725, 29.48 , 32.6 , 37.525, 23.655,
37.8 , 19. , 21.3 , 33.535, 42.46 , 38.95 , 36.1 , 29.3 ,
39.7 , 38.19 , 42.4 , 34.96 , 42.68 , 31.54 , 29.81 , 21.375,
40.81 , 17.4 , 20.3 , 18.5 , 26.125, 41.69 , 24.1 , 36.2 ,
40.185, 39.27 , 34.87 , 44.745, 29.545, 23.54 , 40.47 , 40.66 ,
36.6 , 35.4 , 27.075, 28.405, 21.755, 40.28 , 30.1 , 32.1 ,
23.7 , 35.5 , 29.15 , 27. , 37.905, 22.77 , 22.8 , 34.58 ,
27.1 , 19.475, 26.7 , 34.32 , 24.4 , 41.14 , 22.515, 41.8 ,
26.18 , 42.24 , 26.51 , 35.815, 41.42 , 36.575, 42.94 , 21.01 ,
24.225, 17.67 , 31.5 , 31.1 , 32.78 , 32.45 , 50.38 , 47.6 ,
25.4 , 29.9 , 43.7 , 24.86 , 28.8 , 29.5 , 29.04 , 38.94 ,
44. , 20.045, 40.92 , 35.1 , 29.355, 32.585, 32.34 , 39.8 ,
24.605, 33.99 , 28.2 , 25. , 33.2 , 23.2 , 20.1 , 32.5 ,
37.18 , 46.09 , 39.93 , 35.8 , 31.255, 18.335, 42.9 , 26.79 ,
39.615, 25.9 , 25.745, 28.16 , 23.56 , 40.5 , 35.42 , 39.995,
34.675, 20.52 , 23.275, 36.29 , 32.7 , 19.19 , 20.13 , 23.32 ,
45.32 , 34.6 , 18.715, 21.565, 23. , 37.07 , 52.58 , 42.655,
21.66 , 32. , 18.3 , 47.74 , 22.1 , 19.095, 31.24 , 29.925,
20.35 , 25.85 , 42.75 , 18.6 , 23.87 , 45.9 , 21.5 , 30.305,
44.88 , 41.1 , 40.37 , 28.49 , 33.55 , 40.375, 27.28 , 17.86 ,
33.3 , 39.14 , 21.945, 24.97 , 23.94 , 34.485, 21.8 , 23.3 ,
36.96 , 21.28 , 29.4 , 27.3 , 37.9 , 37.715, 23.76 , 25.52 ,
27.61 , 27.06 , 39.4 , 34.9 , 22. , 30.36 , 27.8 , 53.13 ,
39.71 , 32.87 , 44.7 , 30.97 ])
```

3.Data Visualize : Visualize the unique counts


```
In [12]: 1 sns.countplot(df['bmi'].unique())
```

```
Out[12]: <Axes: ylabel='count'>
```




Find the Null values

```
In [13]: 1 df.isnull().sum()
```

```
Out[13]: age      0  
sex        0  
bmi        0  
children   0  
smoker     0  
region     0  
charges    0  
dtype: int64
```

```
In [14]: 1 Insuranced=df[['age','bmi','children','charges']]  
2 plt.figure(figsize=(4,4))  
3 sns.heatmap(Insuranced.corr(),annot=True)
```

```
Out[14]: <Axes: >
```



To Check The Null values

```
In [15]: 1 df.replace(np.nan, '0', inplace=True)
          2
```

```
In [16]: 1 df.isnull().sum()
```

```
Out[16]: age      0
sex        0
bmi        0
children   0
smoker     0
region     0
charges    0
dtype: int64
```

Feature Scaling: To Split the data into train and test data

```
In [17]: 1 x=np.array(df['age']).reshape(-1,1)
          2 y=np.array(df['charges']).reshape(-1,1)
```

```
In [18]: 1 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
          2 regr=LinearRegression()
          3 regr.fit(x_train,y_train)
          4 print(regr.score(x_test,y_test))
          5
```

0.1124787316394128


In the Linear Regression is not suitable for this model because of accuracy is very less

Logistic Regression

```
In [38]: 1 from sklearn.linear_model import LogisticRegression  
2 from sklearn.preprocessing import StandardScaler
```


```
In [21]: 1 sns.pairplot(df)
```

```
Out[21]: <seaborn.axisgrid.PairGrid at 0x250bfe043d0>
```



```
In [39]: 1 Insuranced=df[['age','bmi']]  
2 plt.figure(figsize=(4,4))  
3 sns.heatmap(Insuranced.corr(),annot=True)
```

Out[39]: <Axes: >



```
In [40]: 1 x = df.iloc[:, :-1].values  
2 y = df.iloc[:, 1].values
```

```
In [41]: 1 #Split the train and test dataset  
2 x_train,x_test,y_train,y_test = train_test_split(x,y,test_size = 0.2)
```

```
In [42]: 1 ml = LogisticRegression()
```

```
In [43]: 1 x=np.array(df['smoker']).reshape(-1,1)
          2 x=np.array(df['age']).reshape(-1,1)
          3 df.dropna(inplace=True)
          4 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=1)
          5 from sklearn.linear_model import LogisticRegression
          6 lr=LogisticRegression(max_iter=10000)
```

```
In [44]: 1 lr.fit(x_train,y_train)
```

Out[44]: LogisticRegression(max_iter=10000)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [45]: 1 score=lr.score(x_test,y_test)
          2 print(score)
```

0.48059701492537316

```
In [46]: 1 sns.scatterplot(data=df,x='smoker',y='charges')
```

```
Out[46]: <Axes: xlabel='smoker', ylabel='charges'>
```



Decesion Tree

```
In [47]: 1 # Decision Tree
2 from sklearn.tree import DecisionTreeClassifier
3 clf=DecisionTreeClassifier()
4 clf.fit(x_train,y_train)
```

```
Out[47]: DecisionTreeClassifier()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [48]: 1 score=clf.score(x_test,y_test)
          2 print(score)
```

0.36716417910447763

Random Forest

```
In [49]: 1 #random forest
          2 from sklearn.ensemble import RandomForestClassifier
          3 rfc=RandomForestClassifier()
          4 rfc.fit(x_train,y_train)
```

Out[49]: RandomForestClassifier()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [50]: 1 params={'max_depth':[2,3,5,10,20],
          2 'min_samples_leaf':[5,10,20,50,100,200],
          3 'n_estimators':[10,25,30,50,100,200]}
```

```
In [51]: 1 from sklearn.model_selection import GridSearchCV
          2 grid_search=GridSearchCV(estimator=rfc,param_grid=params,cv=2,scoring="accuracy")
```

```
In [54]: 1 grid_search.fit(x_train,y_train)
```

Out[54]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
 param_grid={'max_depth': [2, 3, 5, 10, 20],
 'min_samples_leaf': [5, 10, 20, 50, 100, 200],
 'n_estimators': [10, 25, 30, 50, 100, 200]},
 scoring='accuracy')

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.


```
In [55]: 1 grid_search.best_score_
```

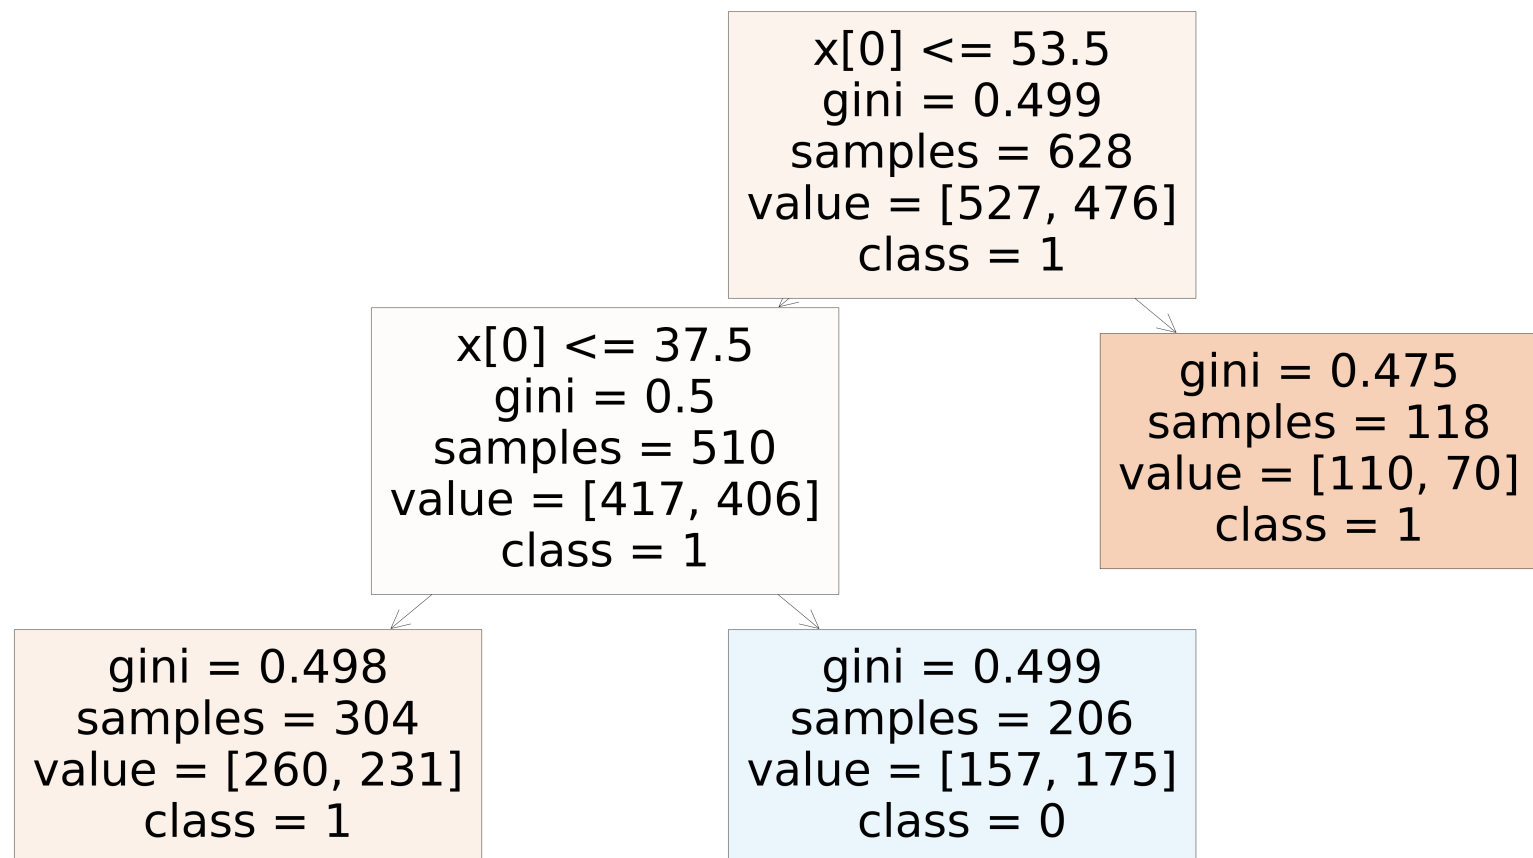
```
Out[55]: 0.5134591375018887
```

```
In [56]: 1 rf_best=grid_search.best_estimator_  
2 rf_best
```

```
Out[56]: RandomForestClassifier(max_depth=2, min_samples_leaf=100, n_estimators=10)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [57]: 1 from sklearn.tree import plot_tree
2 plt.figure(figsize=(80,40))
3 plot_tree(rf_best.estimators_[4],class_names=['1','0'],filled=True);
```



```
In [58]: 1 score=rfc.score(x_test,y_test)
2 print(score)
```

0.36716417910447763

```
In [59]: 1 convert={"sex":{"male":1,'female':2}}
          2 df=df.replace(convert)
          3 df
```

```
Out[59]:
```

| | age | sex | bmi | children | smoker | region | charges |
|------|-----|-----|--------|----------|--------|-----------|-------------|
| 0 | 19 | 2 | 27.900 | 0 | yes | southwest | 16884.92400 |
| 1 | 18 | 1 | 33.770 | 1 | no | southeast | 1725.55230 |
| 2 | 28 | 1 | 33.000 | 3 | no | southeast | 4449.46200 |
| 3 | 33 | 1 | 22.705 | 0 | no | northwest | 21984.47061 |
| 4 | 32 | 1 | 28.880 | 0 | no | northwest | 3866.85520 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 1333 | 50 | 1 | 30.970 | 3 | no | northwest | 10600.54830 |
| 1334 | 18 | 2 | 31.920 | 0 | no | northeast | 2205.98080 |
| 1335 | 18 | 2 | 36.850 | 0 | no | southeast | 1629.83350 |
| 1336 | 21 | 2 | 25.800 | 0 | no | southwest | 2007.94500 |
| 1337 | 61 | 2 | 29.070 | 0 | yes | northwest | 29141.36030 |

1338 rows × 7 columns

```
In [60]: 1 from sklearn.metrics import r2_score
```

```
In [61]: 1 import pickle
```

```
In [62]: 1 filename="Prediction"
          2 pickle.dump(rfc,open(filename,'wb'))
```

Conclusion

for the above different types of models we get accuracy based on the accuracy We can predict the which model is better for this dataset .When we comparing the above accuracies Logistic regression is getting more accuracy among all the models. So, the given dataset is best fit for Logistic Regression

In []:

1